

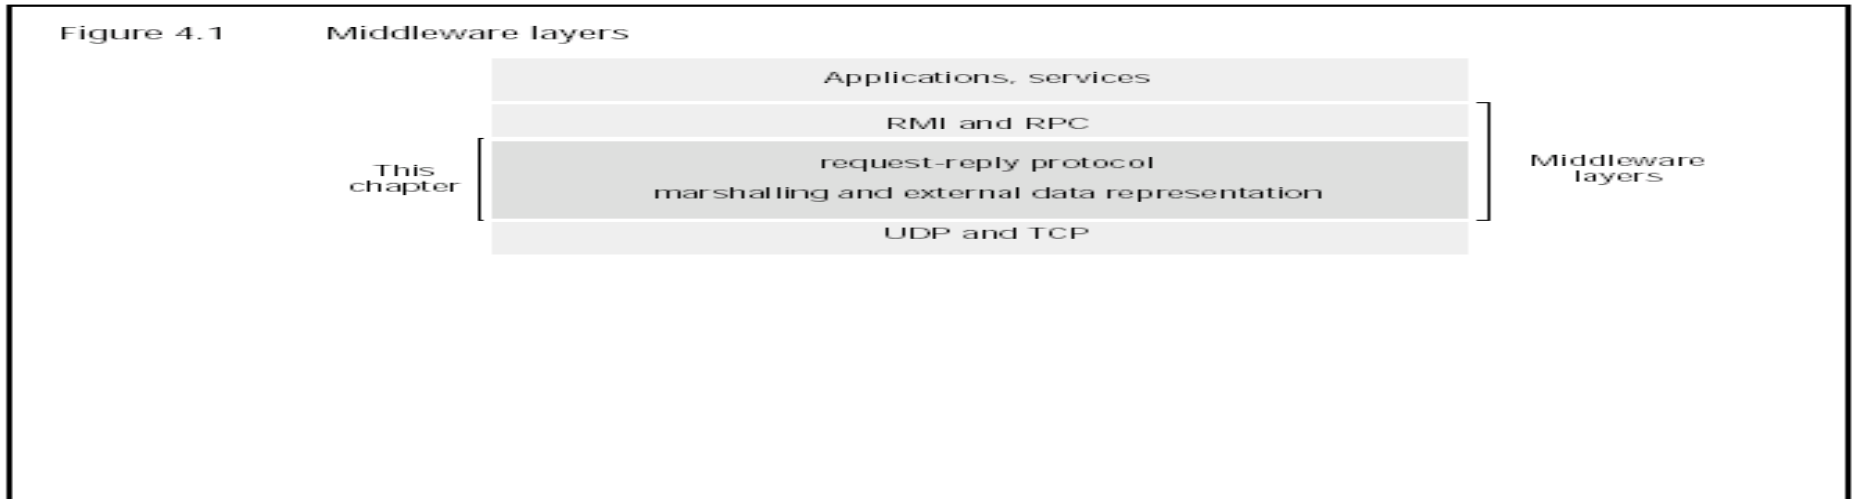
Distributed Systems - Interprocess Communication

- **4. Topics**
 - **4.1 Intro**
 - **4.2 API for Internet Protocols**
 - **4.3 External data representation**
 - **4.4 Client-Server Communication**
 - **4.5 Group communication**
 - **4.6 Unix – An example**

Interprocess Communication – 4.1 Introduction

- Focus:
 - Characteristics of protocols for communication between processes to model distributed computing architecture
 - ♦ Effective means for communicating objects among processes at language level
 - Java API
 - ♦ Provides both datagram and stream communication primitives/interfaces – building blocks for communication protocols
 - Representation of objects
 - ♦ providing a common interface for object references
 - Protocol construction
 - ♦ Two communication patterns for distributed programming: C-S using RMI/RPC and Group communication using 'broadcasting'
 - Unix RPC

Interprocess Communication – 4.1 Introduction



- In **Chapter 3**, we covered Internet transport (TCP/UDP) and network (IP) protocols – without emphasizing how they are used at programming level
- In **Chapter 5**, we cover RMI facilities for accessing remote objects' methods AND the use of RPC for accessing the procedures in a remote server
- **Chapter 4** is on how TCP and UDP are used in a program to effect communication via socket (e.g., Java sockets) – the Middle Layers – for object request/reply invocation and parameter marshalling/representation, including specialized protocols that avoid redundant messaging (e.g., using piggybacked ACKs)

Interprocess Communication – 4.2 API for Internet

- Characteristics of IPC – message passing using send/receive facilities for sync and addressing in distributed programs
- Use of sockets as API for UDP and TCP implementation – much more specification can be found at *java.net*
- Synchronous
 - Queues at remote sites are established for message placement by clients (sender). The local process (at remote site) dequeues the message on arrival
 - If synchronous, both the sender and receiver must ‘rendezvous’ on each message, i.e., both *send* and *receive* invocations are *blocking-until*
- Asynchronous communication
 - *Send* from client is non-blocking and proceeds in parallel with local operations
 - *Receive* could be non-blocking (requiring a background buffer for when message finally arrives, with notification – using interrupts or polling) AND if blocking, perhaps, remote process needs the message, then the process must wait on it
 - Having both sync/async is advantageous, e.g., one thread of a process can do blocked-receive while other thread of same process perform non-block receive or are active – simplifies synchronization. In general non-blocking-receive is simple but complex to implement due to messages arriving out-of-order in the background buffer

Interprocess Communication – 4.2 API for Internet

- Message destinations
 - Typically: send(IP, port#, buffer) – a many-to-one (many senders to a single receiving port), except multicast, which is many-to-group.
 - Possibility: receiving process can have many ports for different message types
 - Server processes usually publish their service-ports for clients
 - Clients can use static IP to access service-ports on servers (limiting, sometimes), but could use location-independent IP by
 - ♦ using name server or binder to bind names to servers at run-time – for relocation
 - ♦ Mapping location-independent identifiers onto lower-level address to deliver/send messages – supporting service migration and relocation
 - IPC can also use 'processes' in lieu of 'ports' for services but ports are flexible and also (a better) support for multicast or delivery to groups of destinations

Interprocess Communication – 4.2 API for Internet


- Reliability
 - Validity: transmission is reliable if packets are delivered despite some drops/losses, and unreliable even if there is a single drop/loss
 - Integrity: message must be delivered uncorrupted and no duplicates
- Ordering
 - Message packets, even if sent out-of-order, must be reordered and delivered otherwise it is a failure of protocol

Interprocess Communication – 4.2 API for Internet

- Sockets
 - Provide an abstraction of endpoints for both TCP and UDP communication
 - Sockets are bound to ports on given computers (via the computer's IP address)
 - Each computer has 2^{16} possible ports available to local processes for receiving messages
 - Each process can designate multiple ports for different message types (but such designated ports can't be shared with other processes on the same computer – unless using IP multicast)
 - Many processes in the same computer can deliver to the same port (many-to-one), however
 - Sockets are typed/associated with either TCP or UDP

Interprocess Communication – 4.2 API for Internet

Figure 4.2 Sockets and ports



Interprocess Communication – 4.2 API for Internet

- Java API for IPs
 - For either TCP or UDP, Java provides an `InetAddress` class, which contains a method: `getByName(DNS)` for obtaining IP addresses, irrespectively of the number of address bits (32 bits for IPv4 or 128 bits for IPv6) by simply passing the DNS hostname. For example, a user Java code invokes:

```
InetAddress aComputer = InetAddress.getByName("nsfcopire.spsu.edu");
```

- The class encapsulates the details of representing the IP address

Interprocess Communication – 4.2 API for Internet

- UDP Datagram communication
 - Steps:
 - ♦ Client finds an available port for UDP connection
 - ♦ Client binds the port to local IP (obtained from `InetAddress.getByName(DNS)`)
 - ♦ Server finds a designated port, publicizes it to clients, and binds it to local IP
 - ♦ Server process issues a receive methods and gets the IP and port # of sender (client) along with the message
 - Issues
 - ♦ **Message size** – set to 8KByte for most, general protocol support 216 bytes, possible truncation if receiver buffer is smaller than message size
 - ♦ **Blocking** – send is non-blocking and `op` returns if message gets pass the UDP and IP layers; receive is blocking (with discard if no socket is bound or no *thread* is waiting at destination port)
 - ♦ **Timeouts** – reasonably large time interval set on receiver sockets to avoid indefinite blocking
 - ♦ **Receive from any** – no specification of sources (senders), typically many-to-one, but one-to-one is possible by a designated send-receive socket (known by both C/S)

Interprocess Communication – 4.2 API for Internet

- UDP Failure Models:
 - Due to Omission of send or receive (either checksum error or no buffer space at source or destination)
 - Due to out-of-order delivery
 - UDP lacks built in checks, but failure can be modeled by implementing an ACK mechanism

Interprocess Communication – 4.2 API for Internet

- Use of UDP – Client/Sender code

Figure 4.3 UDP client sends a message to the server and gets a reply

```
import java.net.*;
import java.io.*;
public class UDPClient{
    public static void main(String args[]) {
        // args give message contents and server hostname
        DatagramSocket aSocket = null;
        try {
            aSocket = new DatagramSocket();
            byte [] m = args[0].getBytes();
            InetAddress aHost = InetAddress.getByName(args[1]);
            int serverPort = 6789;
            DatagramPacket request = new DatagramPacket(m, args[0].length(), aHost, serverPort);
            aSocket.send(request);
            byte[] buffer = new byte[1000];
            DatagramPacket reply = new DatagramPacket(buffer, buffer.length);
            aSocket.receive(reply);
            System.out.println("Reply: " + new String(reply.getData()));
        } catch (SocketException e) {System.out.println("Socket: " + e.getMessage());}
        } catch (IOException e) {System.out.println("IO: " + e.getMessage());}
        } finally { if(aSocket != null) aSocket.close();}
    }
}
```

Interprocess Communication – 4.2 API for Internet

- Use of UDP – Server/Receiver code

Figure 4.4 UDP server repeatedly receives a request and sends it back to the client

```
import java.net.*;
import java.io.*;
public class UDPServer{
    public static void main(String args[]){
        DatagramSocket aSocket = null;
        try{
            aSocket = new DatagramSocket(6789);
            byte[] buffer = new byte[1000];
            while(true){
                DatagramPacket request = new DatagramPacket(buffer, buffer.length);
                aSocket.receive(request);
                DatagramPacket reply = new DatagramPacket(request.getData(),
                    request.getLength(), request.getAddress(), request.getPort());
                aSocket.send(reply);
            }
        } catch (SocketException e) {System.out.println("Socket: " + e.getMessage());}
        } catch (IOException e) {System.out.println("IO: " + e.getMessage());}
        } finally {if(aSocket != null) aSocket.close();}
    }
}
```

Interprocess Communication – 4.2 API for Internet

- TCP Stream Communication

- Grounded in the ‘piping’ architecture of Unix systems using BSD Unix sockets for streaming bytes
- Characteristics:
 - ♦ Message sizes – user application has option to set IP packet size, small or large
 - ♦ Lost messages – Sliding window protocol with ACKs and retransmission is used
 - ♦ Flow control – Blocking or throttling is used
 - ♦ Message duplication and ordering – Seq #s with discard of dups & reordering
 - ♦ Message destinations – a connection is established first, using connection-accept methods for rendezvous, and no IP addresses in packets. [Each connection socket is bidirectional – using two streams: output/write and input/read]. A client closes a socket to sign off, and last stream of bytes are sent to receiver with ‘broken-pipe’ or empty-queue indicator

Interprocess Communication – 4.2 API for Internet

- TCP Stream Communication

- Other Issues

- ♦ Matching of data items – both client/sender and server/receiver must agree on data types and order in the stream
 - ♦ Blocking – data is streamed and kept in server queue: empty server queue causes a block AND full server queue causes a blocking of sender
 - ♦ Threads – used by servers (in the background) to service clients, allowing asynchronous blocking. [Systems without threads, e.g., Unix, use *select*]

- Failure Model

- ♦ Integrity: uses checksums for detection/rejection of corrupt data and seq #s for rejecting duplicates
 - ♦ Validity: uses timeout with retransmission techniques (takes care of packet losses or drops)
 - ♦ Pathological: excessive drops/timeouts signal broken sockets and TCP throws in the towel (no one knows if pending packets were exchanged) – unreliable

- Uses – TCP sockets used for such services as: HTTP, FTP, Telnet, SMTP

Interprocess Communication – 4.2 API for Internet

- Use of TCP – Client/Sender code

Figure 4.5 TCP client makes connection to server, sends request and receives reply

```
import java.net.*;
import java.io.*;
public class TCPClient {
    public static void main (String args[]) {
        // arguments supply message and hostname of destination
        Socket s = null;
        try{
            int serverPort = 7896;
            s = new Socket(args[1], serverPort);
            DataInputStream in = new DataInputStream( s.getInputStream());
            DataOutputStream out =
                new DataOutputStream( s.getOutputStream());
            out.writeUTF(args[0]);    // UTF is a string encoding see Sn 4.3
            String data = in.readUTF();
            System.out.println("Received: " + data) ;
        } catch (UnknownHostException e) {
            System.out.println("Sock: " + e.getMessage());
        } catch (EOFException e) {System.out.println("EOF: " + e.getMessage());}
        } catch (IOException e) {System.out.println("IO: " + e.getMessage());}
    } finally {if(s!=null) try {s.close();} catch (IOException e) {/*close failed*/}}
    }
}
```


Interprocess Communication – 4.2 API for Internet

- Use of TCP – Server/Receiver code

Figure 4.6 TCP server makes a connection for each client and then echoes the client's request

```
import java.net.*;
import java.io.*;
public class TCPServer {
    public static void main (String args[]) {
        try{
            int serverPort = 7896;
            ServerSocket listenSocket = new ServerSocket(serverPort);
            while(true) {
                Socket clientSocket = listenSocket.accept();
                Connection c = new Connection(clientSocket);
            }
        } catch(IOException e) {System.out.println("Listen :"+e.getMessage());}
    }
}

// this figure continues on the next slide
```

Interprocess Communication – 4.2 API for Internet

Use of TCP – Server/Receiver code (cont'd)

Figure 4.6 continued

```
class Connection extends Thread {
    DataInputStream in;
    DataOutputStream out;
    Socket clientSocket;
    public Connection (Socket aClientSocket) {
        try {
            clientSocket = aClientSocket;
            in = new DataInputStream( clientSocket.getInputStream());
            out =new DataOutputStream( clientSocket.getOutputStream());
            this.start();
        } catch(IOException e) {System.out.println("Connection:"+e.getMessage());}
    }
    public void run(){
        try {                                // an echo server
            String data = in.readUTF();
            out.writeUTF(data);
        } catch EOFException e) {System.out.println("EOF:"+e.getMessage());}
        } catch(IOException e) {System.out.println("IO:"+e.getMessage());}
        } finally { try {clientSocket.close();} catch (IOException e){/*close failed*/}}
    }
}
```

Interprocess Communication – 4.3 External data representation

- Issues
 - At language-level data (for comm) are stored in data structures
 - At TCP/UDP-level data are communicated as ‘messages’ or streams of bytes – hence, conversion/flattening is needed
 - Problem? Different machines have different primitive data reps, e.g., big-endian and little-endian order of integers, float-type, char codes
 - Marshalling (before trans) and unmarshalling (restored to original on arrival)
 - Either both machines agree on a format type (included in parameter list) or an *intermediate* external standard (external data rep) is used, e.g., CORBA Common Data Rep (CDR)/IDL for many languages; Java object serialization for Java code only, Sun XDR standard for Sun NFSs

Interprocess Communication – 4.3 External data representation

- This masks the differences due to different computer hardware.
- CORBA CDR
 - only defined in CORBA 2.0 in 1998, before that, each implementation of CORBA had an external data representation, but they could not generally work with one another. That is:
 - ♦ the heterogeneity of hardware was masked
 - ♦ but not the heterogeneity due to different programmers (until CORBA 2)
 - CORBA CDR represents simple and constructed data types (sequence, string, array, struct, enum and union)
 - ♦ note that it does not deal with objects (*only Java does: objects and tree of objects*)
 - it requires an IDL specification of data to be serialised
- Java object serialisation
 - represents both objects and primitive data values
 - it uses reflection to serialise and deserialise objects– it does not need an IDL specification of the objects. (Reflection: inquiring about class properties, e.g., names, types of methods and variables, of objects]

Interprocess Communication – 4.3 External data representation

- **Example of Java serialized message**

```
public class Person implements Serializable {  
    private String name;  
    private String place;  
    private int year;  
    public Person(String aName, String aPlace, int aYear) {  
        name = aName;  
        place = aPlace;  
        year = aYear;  
    }  
    // followed by methods for accessing the instance variables  
}
```

– Consider the following object:

Person p = new Person("Smith", "London", 1934);

Figure 4.9 Indication of Java serialized form

Serialized values			Explanation
Person	8-byte version number	h0	class name, version number
3	int year	java.lang.String name:	number, type and name of instance variables
1934	5 Smith	6 London	values of instance variables
		h1	

The true serialized form contains additional type markers; h0 and h1 are handles

CORBA IDL example

CORBA has a struct

```
struct Person {  
    string name;  
    string place;  
    long year;
```

```
};
```

```
interface PersonList {  
    readonly attribute string listname;  
    void addPerson(in Person p) ;  
    void getPerson(in string name, out Person
```

```
p);
```

```
    long number();
```

```
};
```


remote interface

remote interface defines
methods for RMI

parameters are *in*, *out* or *inout*

- Remote interface:
 - specifies the **methods** of an object available for remote invocation
 - an interface definition language (or IDL) is used to specify remote interfaces. E.g. the above in CORBA IDL.
 - Java RMI would have a class for *Person*, but CORBA has a *struct*


Interprocess Communication – 4.3 External data representation



- each process contains objects, some of which can receive remote invocations, others only local invocations
- those that can receive remote invocations are called *remote objects*
- objects need to know the *remote object reference* of an object in another process in order to invoke its methods. **How do they get it?**
- the *remote interface* specifies which methods can be invoked remotely
- Remote object references are passed as arguments and compared to ensure uniqueness over time and space in Distributed Computing system

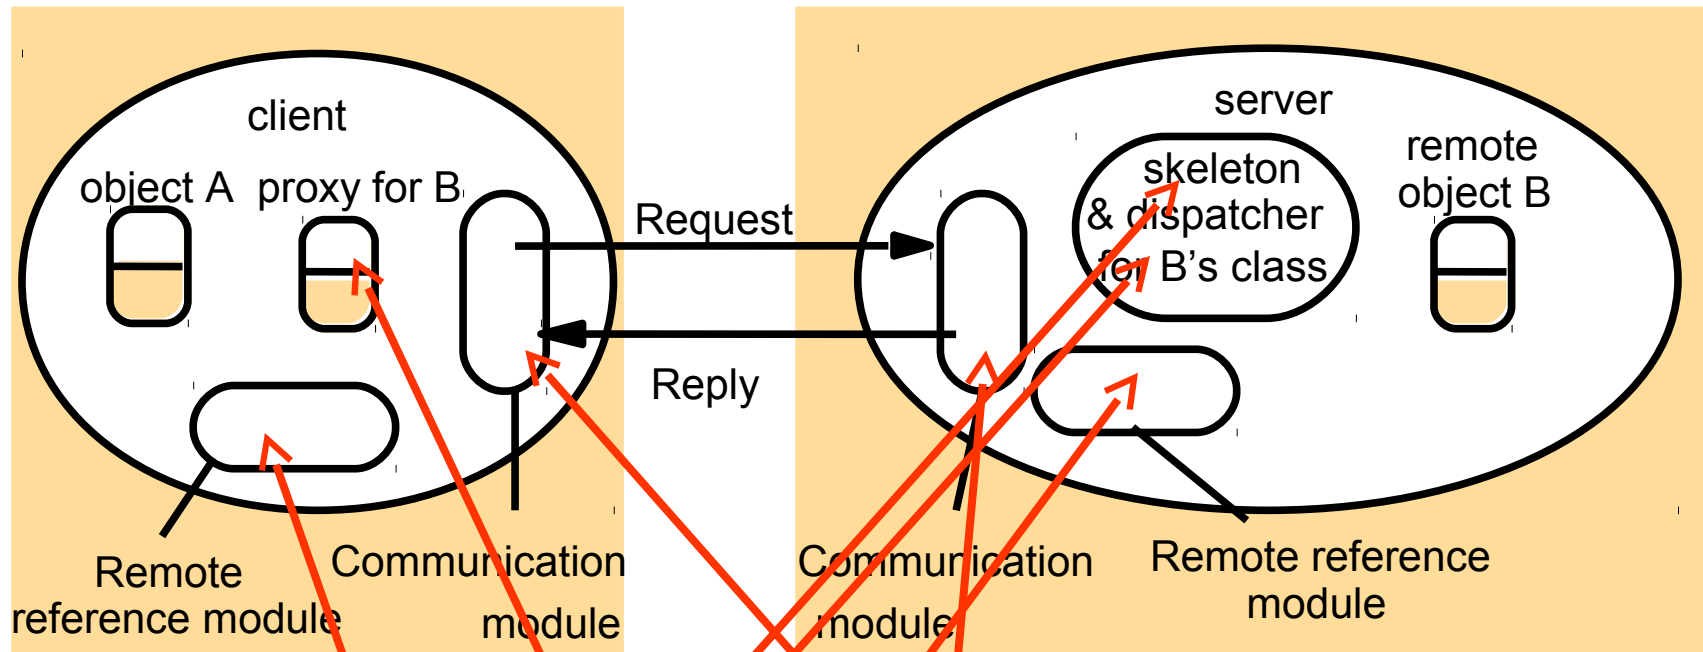
Representation of a remote object reference

Figure 4.10



- a remote object reference must be unique in the distributed system and over time. It should not be reused after the object is deleted. **Why not?**
- the first two fields locate the object unless migration or re-activation in a new process can happen
- the fourth field identifies the object within the process
- its interface tells the receiver what methods it has (e.g. class *Method*)
- a remote object reference is created by a remote reference module when a reference is passed as argument or result to another process
 - it will be stored in the corresponding proxy
 - it will be passed in request messages to identify the remote object whose method is to be invoked

The architecture of remote method invocation



Proxy - makes RMI transparent to remote interface. Marshals request results. Forwards request.

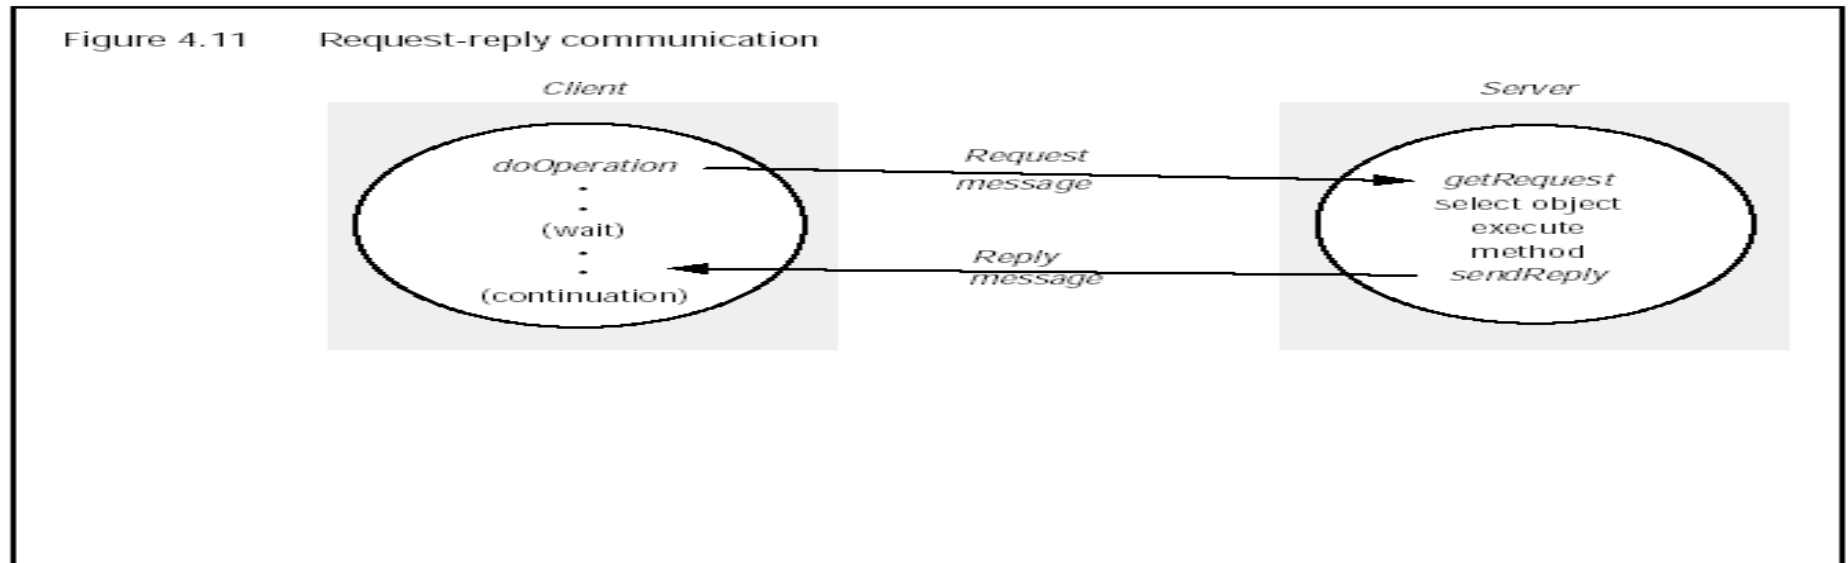
carries out Request-reply protocol

translates between local and remote object references and creates remote object references. Uses remote object table

RMI software - between application level objects and communication and remote reference modules

Interprocess Communication – 4.4 Client-Server Communication

- Modes:
 - Request-reply: client process blocks until and ACK is received from server (Synchronous)
 - Use send/receive operations in Java API for UDP (or TCP streams – typically with much overhead for the ‘guarantees’)
 - Protocol over UDP, e.g., piggybacked ACKs,



Interprocess Communication – 4.4 Client-Server Communication

Figure 4.12 Operations of the request-reply protocol

```
public byte[] doOperation (RemoteObjectRef o, int methodId, byte[] arguments)  
    sends a request message to the remote object and returns the reply.  
    The arguments specify the remote object, the method to be invoked and the  
    arguments of that method.  
  
public byte[] getRequest ();  
    acquires a client request via the server port.  
  
public void sendReply (byte[] reply, InetAddress clientHost, int clientPort);  
    sends the reply message reply to the client at its Internet address and port.
```

Request-Reply Protocol

Figure 4.13 Request-reply message structure

messageType
requestId
objectReference
methodId
arguments

int (0=Request, 1=Reply)
int
RemoteObjectRef
int or *Method*
// array of bytes

MessageIDs: *requestID + IP.portnumber* // IP.portnumber from packet if UDP

Interprocess Communication – 4.4 Client-Server Communication

- Failure Model of Request-Reply Protocol
 - For doOperation, getRequest, sendReply over UDP, the possible problems include:
 - ♦ Omission failure (link failures, drops/losses, missed/corrupt addresses)
 - ♦ Out-of-order delivery
 - ♦ Node/process down
 - Solved by
 - ♦ Timeouts with retrans or retries until reply is received/confirmed
 - ♦ Discards of repeated requests (by server process)
 - ♦ On lost reply messages, servers repeats idempotent operations
 - ♦ Maintain history (reqid, message, client-id) or buffer replies and retrans – memory intensive

Interprocess Communication – 4.4 Client-Server Communication

- Failure handling –
 - ♦ RPC exchange protocols (for handling failures)
 - Request (R) Protocol – for remote procedures with return results or need for ACK
 - Request-Reply Protocol – converse of R protocol, with reply message as piggybacked ACK
 - Request-Reply-ACK-reply (RRA) – request and ACK reply message (with id of specific range of requests being ACKed – Go-back/Selective Repeat protocols) – avoids keeping histories

Figure 4.14 RPC exchange protocols

<i>Name</i>	<i>Messages sent by</i>		
	<i>Client</i>	<i>Server</i>	<i>Client</i>
R	<i>Request</i>		
RR	<i>Request</i>	<i>Reply</i>	
RRA	<i>Request</i>	<i>Reply</i>	<i>Acknowledge reply</i>

Interprocess Communication – 4.5 Group Communication

- Limitations of peer-to-peer (point-to-point)
- Support for (concurrent) multiple services from a single clients requires a different communication paradigm
- Multicast – message from a single client to group of server processes (useful for distributed processing)
- Characteristics
 - Fault tolerance based on replicated services (redundancy of same service)
 - Finding discovery servers (with registry of new service interfaces)
 - Better performance via data replication (updates are multicast for currency)
 - Propagation of event notification (inform affected servers of new events, e.g., new routing tables, new servers or clients)

Interprocess Communication – 4.5 Group Communication

- IP Multicast

- Built on top of the IP layer (using IP addresses) [ports for TCP and UDP]
- Thus, multicast packets are sent to computers using an IP-addressing scheme
- Client/sender is unaware of individual computers in a group
- A multicast-group address is specified using class D Internet address
- Membership of a group is dynamic
- Multicasting is through UDP protocol only (at programming level) via multicast addresses and ordinary port numbers
- A receiver joins a group by adding its socket to the group
- At IP level, adding process's socket #s to a group makes the computer a member of the group – allowing all such sockets to receive incoming packets at specified port #s.
- How? Via local multicast capability Or via multicast routers @ Internet level
- Membership? Permanent (designated addresses) or temporary

Interprocess Communication – 4.5 Group Communication

IP Multicast – Example @ Application level

Figure 4.17 Multicast peer joins a group and sends and receives datagrams

```
import java.net.*;
import java.io.*;
public class MulticastPeer{
    public static void main(String args[]){
        // args give message contents & destination multicast group (e.g. "228.5.6.7")
        MulticastSocket s =null;
        try {
            InetAddress group = InetAddress.getByName(args[1]);
            s = new MulticastSocket(6789);
            s.joinGroup(group);
            byte [] m = args[0].getBytes();
            DatagramPacket messageOut =
                new DatagramPacket(m, m.length, group, 6789);
            s.send(messageOut);

            // this figure continued on the next slide
        } catch (Exception e){
            e.printStackTrace();
        }
    }
}
```


Interprocess Communication – 4.5 Group Communication

Figure 4.17 continued

```
// get messages from others in group
byte[] buffer = new byte[1000];
for(int i=0; i< 3; i++) {
    DatagramPacket messageIn =
        new DatagramPacket(buffer, buffer.length);
    s.receive(messageIn);
    System.out.println("Received:" + new String(messageIn.getData()));
}
s.leaveGroup(group);
} catch (SocketException e){System.out.println("Socket: " + e.getMessage());}
} catch (IOException e){System.out.println("IO: " + e.getMessage());}
} finally { if(s != null) s.close();}
}
```

Interprocess Communication – 4.6 UNIX Example

Figure 4.18 Sockets used for datagrams

Sending a message

```
s = socket(AF_INET, SOCK_DGRAM, 0)
.  
.  
bind(s, ClientAddress)
.  
.  
sendto(s, "message", ServerAddress)
```

Receiving a message

```
s = socket(AF_INET, SOCK_DGRAM, 0)
.  
.  
bind(s, ServerAddress)
.  
.  
amount = recvfrom(s, buffer, from)
```

ServerAddress and *ClientAddress* are socket addresses

Interprocess Communication – 4.6 UNIX Example

Figure 4.19 Sockets used for streams

Requesting a connection

```
s = socket(AF_INET, SOCK_STREAM, 0)
.  
.  
connect(s, ServerAddress)
.  
.  
write(s, "message", length)
```

Listening and accepting a connection

```
s = socket(AF_INET, SOCK_STREAM, 0)
.  
bind(s, ServerAddress);  
listen(s, 5);  
.  
sNew = accept(s, ClientAddress);  
.  
n = read(sNew, buffer, amount)
```

ServerAddress and *ClientAddress* are socket addresses