

# Network File Systems (NFS): Case Study

- **NFS (Network File System)**
  - Developed by Sun Microsystems (in 1985)
  - Most popular, open, and widely used.
  - NFS protocol standardized through IETF (RFC 1813)

# NFS Architecture

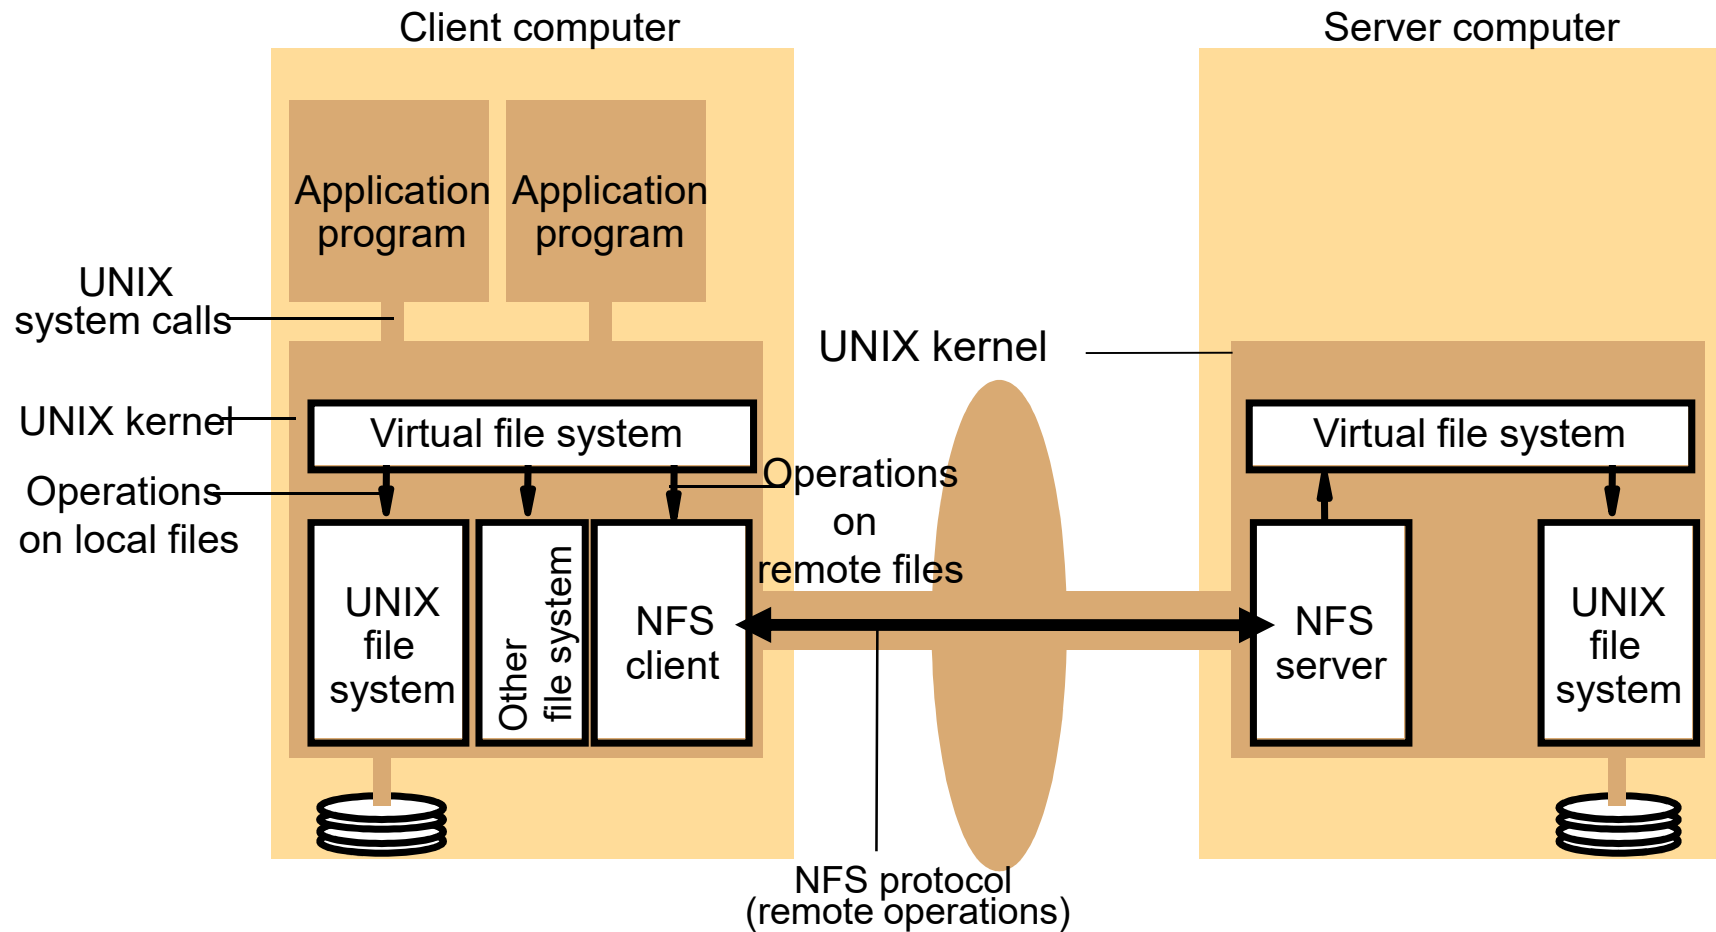



Figure: NFS architecture

# Client side (III)

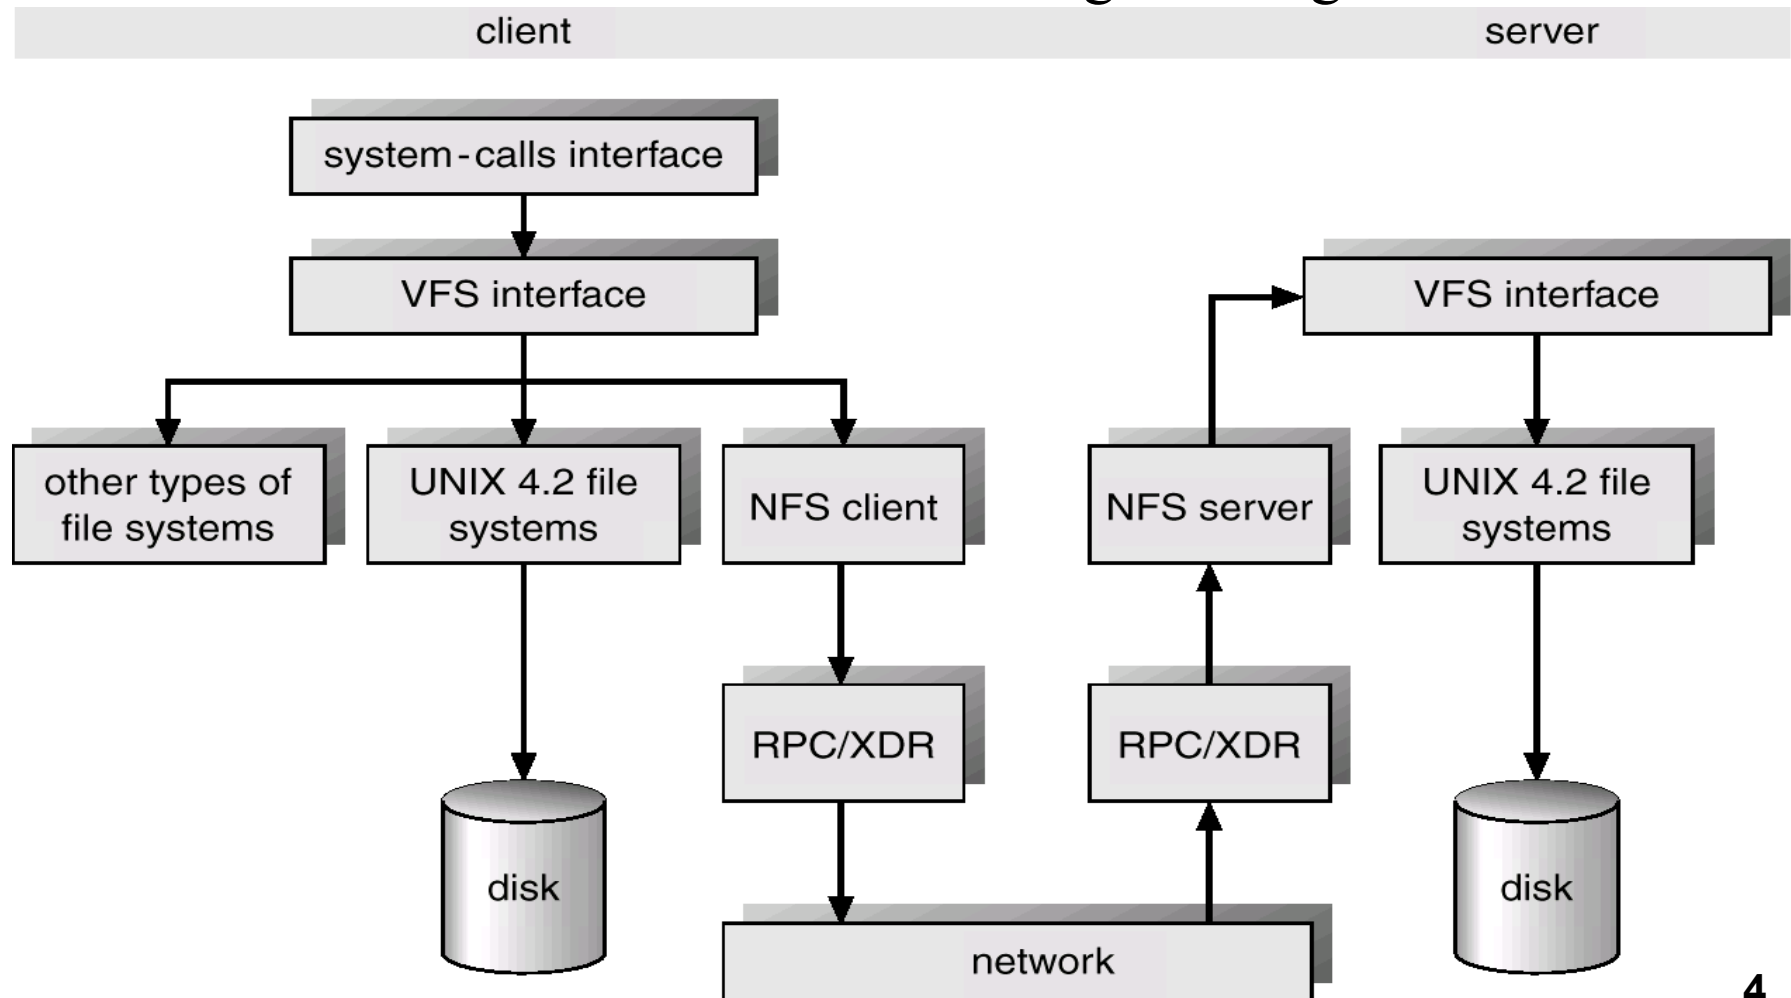


# DISTRIBUTED FILE SYSTEMS

## SUN Network File System

### NFS ARCHITECTURE:

Follow local and remote access through this figure:



# Sun NFS

- The file identifiers used in NFS are called file handles.

fh = file handle:

Filesystem identifier

i-node number

i-node generation

# Server side

- **File handle** consists of
  - **Filesystem id** identifying disk partition
  - **I-node number** identifying file within partition
  - **Generation number** changed every time. i-node number is i-node number is reused after a file is removed (to store a new file)
  - **V-node-** contains an indicator to show whether a file is local or remote. If file is local, the v-node contain reference to the index of local file(i-node). If the file system is remote, it contains the file handle to remote file.
- **Filesystem id** in filesystem **superblock**(File system type, Size, Status, Information about other metadata structures)

# **DISTRIBUTED FILE SYSTEMS**

## **SUN Network File System**

### **NFS ARCHITECTURE:**

1. UNIX filesystem layer - does normal open / read / etc. commands.
2. Virtual file system ( VFS ) layer –
  - a) Gives clean layer between user and filesystem.
  - b) Acts as deflection point by using global vnodes.
  - c) Understands the difference between local and remote names.
  - d) Keeps in memory information about what should be deflected (mounted directories) and how to get to these remote directories.
3. System call interface layer -
  - a) Presents sanitized validated requests in a uniform way to the VFS.

# Sun NFS

- Mount service

- Mount operation:

- `mount(remotehost, remotedirectory, localdirectory)`

- Server maintains a table of clients who have mounted filesystems at that server.

- Each client maintains a table of mounted file systems holding:

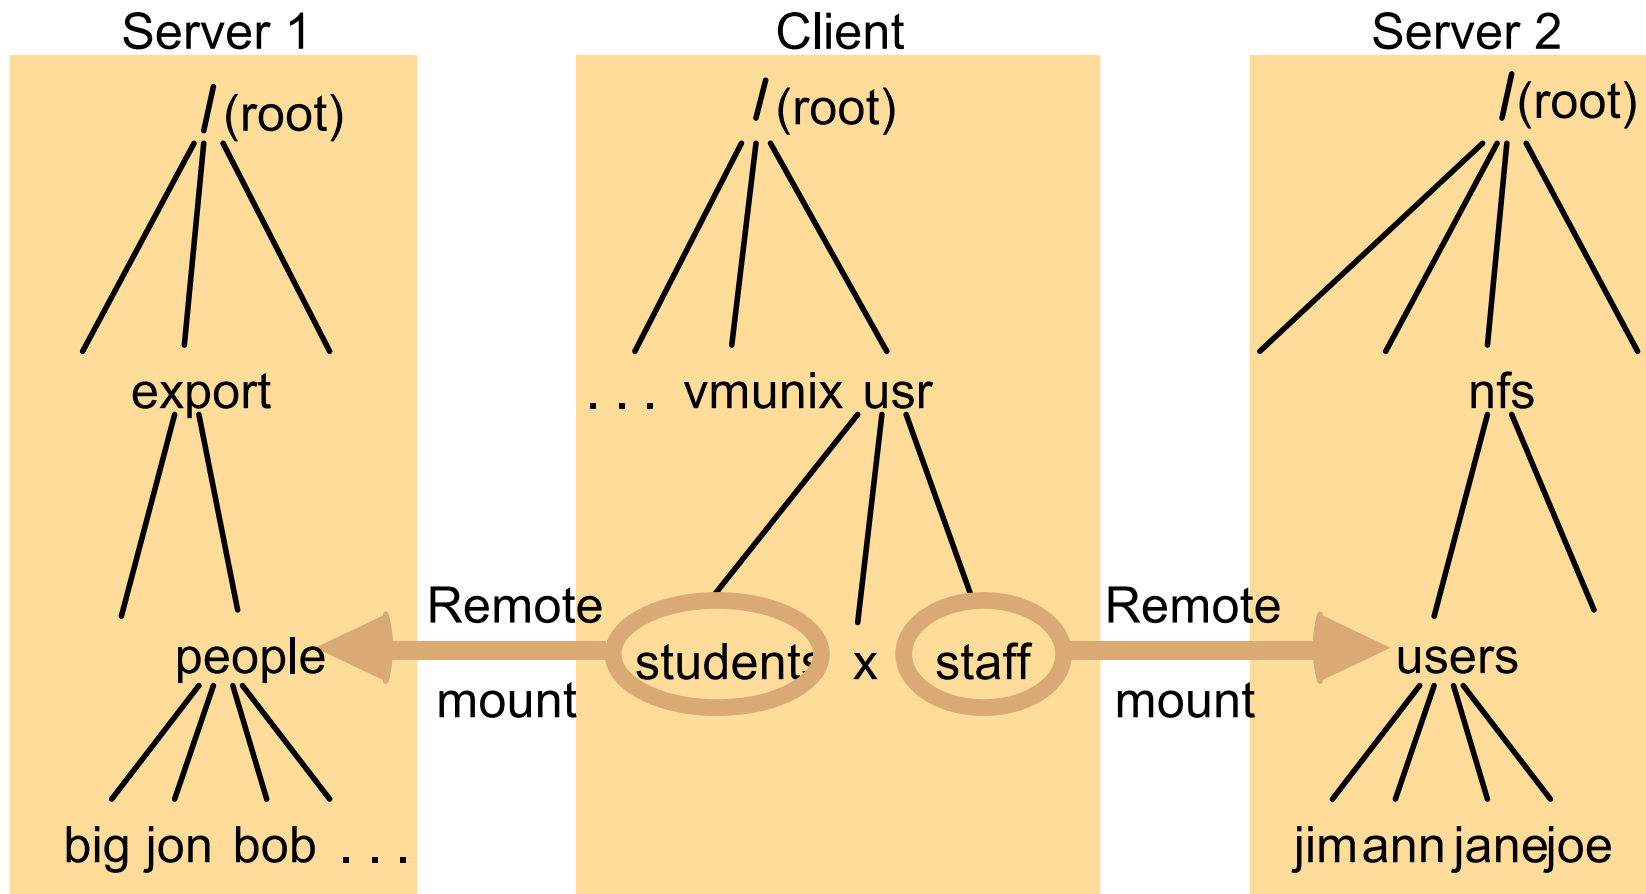
- `< IP address, port number, file handle >`

- Remote file systems may be hard-mounted (In case of failure, it will repeatedly retry to connect the server) or soft-mounted (In case of failure, it will report an error) in a client computer.

- Figure 10 illustrates a Client with two remotely mounted file stores.



# Sun NFS



**Figure : Local and remote file systems accessible on an NFS client**

Note: The file system mounted at `/usr/students` in the client is actually the sub-tree located at `/export/people` in Server 1; the file system mounted at `/usr/staff` in the client is actually the sub-tree located at `/nfs/users` in Server 2.

# Sun NFS

- NFS summary
  - NFS is an excellent example of a simple, robust, high-performance distributed service.
  - Achievement of transparencies are other goals of NFS:
    - ❖ Access transparency:
      - Enables local and remote resources to be accessed using identical operations.
      - The API is the UNIX system call interface for both local and remote files.
    - ❖ Location transparency:
      - NFS provides the location transparency i.e name does not hint at its physical storage location.

# Sun NFS

- ❖ Mobility transparency:
  - Hardly achieved; relocation of filesystems is possible, but requires updates to client configurations.
- ❖ Scalability transparency:
  - NFS does not scale well as it is limited to LAN
- ❖ Replication transparency:
  - Limited to read-only file systems; for writable files, the SUN Network Information Service (NIS) runs over NFS and is used to replicate essential system files.
- ❖ Hardware and software operating system heterogeneity:
  - NFS has been implemented for almost every known operating system and hardware platform and is supported by a variety of filling systems.

# Sun NFS

- ❖ Fault tolerance:
  - Limited but effective; service is suspended if a server fails. Recovery from failures is aided by the simple stateless design.

# Sun NFS

## CACHES OF REMOTE DATA:

- The client keeps:
  - File block cache - (the contents of a file)
  - File attribute cache - (file header info (inode in UNIX)).
- The local kernel hangs on to the data after getting it the first time.
- On an open, local kernel, it checks with server that cached data is still OK.
- Cached attributes are thrown away after a few seconds.